



I'm building a t-shirt/merch design generator using AI image generation APIs. I need current (December 2025) information. Please only cite sources from 2025, prioritizing the most recent data available:

1. Current OpenAI Image Models

- What image generation models does OpenAI currently offer as of December 2025?
- What are the differences between them (quality, speed, text rendering, price)?
- Which is recommended for commercial graphic design with text?
- What's the latest/newest model and what are its capabilities?

2. API vs ChatGPT

- Is there still a quality difference between using OpenAI's image API directly vs ChatGPT's interface?
- If so, what causes this gap and can it be worked around?
- Does the API auto-enhance/rewrite prompts? Can this be controlled?

3. Text/Typography in AI Images

- Which current AI image model (OpenAI or otherwise) has the BEST text rendering accuracy?
- Can font styles be controlled, or do models default to the same few fonts?
- What prompt techniques produce accurate AND varied typography?
- What's the maximum reliable text length?

4. Optimal Prompt Structure

- What is the current best practice for structuring image generation prompts?
- What elements MUST be included for best results?
- What should be AVOIDED?
- Are there style/quality parameters that affect output significantly?

5. T-Shirt/Merch Design Specific

- Best prompt techniques for print-on-demand apparel graphics?
- How to reliably get transparent or solid color backgrounds?
- Recommended size/aspect ratio for apparel?
- Known issues with print-ready design generation?

6. Model Comparison (All Providers)

- As of December 2025, what are the top AI image generators for text-heavy commercial designs?
- Compare: OpenAI models, Flux, Ideogram, Midjourney, Google Imagen, and any other leaders
- Which offers the best typography control?

7. API Implementation

- Should I use an LLM to pre-process/enhance prompts before sending to the image API?
- Any prompt caching or optimization techniques?
- What's the recommended workflow for consistent, high-quality results?

Please provide specific examples and cite sources with dates.

OpenAI's current primary image model is GPT-Image-1 (plus cheaper "mini"-style variants), and it is competitive but not state-of-the-art for complex typography compared with specialists like Ideogram 2/3 and Midjourney's latest versions. For a t-shirt/merch generator with heavy text, the practical stack in late 2025 is: Ideogram (2.0/3.0) or Midjourney v7 for text-centric layouts, GPT-Image-1 if you want tight integration with your LLM/backend, and Flux.2 / Imagen 4 as strong alternates depending on your infra. [1] [2] [3] [4] [5] [6] [7] [8]

1. Current OpenAI image models

- OpenAI's public image lineup in Dec 2025 centers on **gpt-image-1**, described as its "state-of-the-art image generation model," exposed both via /v1/images/generations and via chat/assistants. [4] [9] [10]
- **Capabilities:** multimodal (text+image in, image out), 1024×1024 and wide formats (e.g. 1536×1024), strong prompt adherence, decent text rendering, and the same trust/safety stack and C2PA watermarking as ChatGPT's image tool. [11] [12] [4]

Quality / speed / price (2025):

- GPT-Image-1 is positioned as "higher" quality but "slowest" compared with cheaper image variants, with pricing on the order of **≈\$0.07 per medium image** at standard quality, billed in "image output tokens" (around \$40 per million image tokens). [4] [11]
- Third-party aggregators describe **gpt-image-1-mini** (or similar names) as a **cost-optimized** variant with lower quality and faster speed, useful for thumbnails or drafts. [13]

Best OpenAI model for commercial graphics with text:

- For OpenAI specifically, **GPT-Image-1** is the recommended choice for merch designs because it has the best prompt alignment and text rendering within OpenAI's stack, and it's explicitly built for professional image generation via API.[\[14\]](#) [\[11\]](#) [\[4\]](#)

Latest / newest model & roadmap nuance:

- GPT-Image-1 is currently documented as the latest stable image model in the API docs as of late Nov 2025.[\[9\]](#) [\[14\]](#) [\[4\]](#)
- TestingCatalog and other watchers report OpenAI is **testing new "Image-2" and "Image-2-mini" models** on LM Arena (Dec 9, 2025), claiming higher detail and color accuracy in ChatGPT; this suggests a near-term successor but it is not yet a documented production API model.[\[15\]](#)

2. API vs ChatGPT for image quality

- OpenAI's dev forum announcement in April 2025 confirms that **the same image model powering ChatGPT is now available in the API**, which removes the historic quality gap where ChatGPT sometimes used a newer internal model.[\[10\]](#) [\[4\]](#)
- Any residual differences you see now are mostly about:
 - **Prompt rewriting:** ChatGPT applies an LLM "prompt engineer" layer; the image API does not, unless you build that yourself.
 - **System safety transforms:** The same policy stack applies (content filters, watermarking), but ChatGPT UX may apply extra guardrails or style normalizations.

Can the gap be worked around?

- You can **replicate ChatGPT-style prompting** by using a GPT-4o/5.x-class model to rewrite users' design intents into a rich, structured image prompt, then passing that to /v1/images/generations with model: "gpt-image-1". This gives you the same underlying model while controlling the pre-processing yourself.[\[14\]](#) [\[4\]](#)

Prompt auto-enhancement in the API:

- The **image API itself does not publicly expose an "auto-enhance" toggle**; it accepts your prompt verbatim, subject only to safety filtering.[\[11\]](#) [\[4\]](#)
- Any "smart" reformulation has to be done by your own LLM layer, which you can fully control (e.g., strict JSON schema for layout, color, fonts).

3. Text / typography in AI images (2025)

Who is best for text now?

Across 2025 reviews and vendor claims, the leading text-heavy generators are:

- **Ideogram 2.0 / 3.0** – explicitly marketed as having "industry-leading" text rendering accuracy and premium graphic design capabilities, including long, stylized text; the

"Design" style is tuned specifically for typography in greeting cards, posters, and POD graphics.[\[5\]](#) [\[7\]](#) [\[8\]](#) [\[16\]](#)

- **Midjourney v7** – reviewer testing ranks it among the best for accurate, styled text; v7 adds "dedicated text rendering capabilities" that make it viable for professional text-heavy work.[\[6\]](#)
- **Google Imagen 4 (Vertex AI)** – Google Cloud positions Imagen 4 as having "**outstanding text rendering and prompt adherence**" for enterprise workflows.[\[2\]](#)
- **GPT-Image-1** – CometAPI notes that GPT-Image-1 is particularly good for **crisp text blocks** in infographics/diagrams, but most independent 2025 roundups still place it slightly behind Ideogram/MJ/Imagen for long decorative typography.[\[17\]](#) [\[18\]](#) [\[6\]](#)

Font control vs "mushy" defaults:

- None of the major models provide full font-file-level control (e.g., "exactly set Futura PT 700"), but:
 - **Ideogram 2.x/3.0** exposes "Design" style and typography-specific controls; its prompt guide emphasizes describing **font family category** (serif/sans/handwritten/blackletter) and layout (arched, stacked, badge, logo), which it follows more consistently than most models.[\[7\]](#) [\[8\]](#) [\[16\]](#)
 - **Midjourney v6/v7** supports text by putting the target string in quotes and describing font style ("bold condensed sans-serif", "retro script"), plus tuning --stylize and --style raw for clarity vs artistry; tests show it varies font style well but not with precise brand-font fidelity.[\[19\]](#) [\[20\]](#) [\[21\]](#) [\[6\]](#)
 - **Imagen 3/4** lets you specify decorative text in greeting cards, posters etc., with control over "various fonts and colors," but Google itself admits some remaining weakness with complex layouts and scripts.[\[22\]](#) [\[23\]](#) [\[2\]](#)

Effective prompt techniques for accurate + varied typography:

From the 2025 typography guides (Ideogram and Midjourney) and benchmarking articles:[\[23\]](#)
[\[16\]](#) [\[19\]](#) [\[6\]](#)

- Put the exact text **in quotes**, and keep it short and clean:
 - Example: t-shirt graphic, center badge logo reading "NEON TIGER", bold condensed sans-serif, all caps, high contrast, no extra text.[\[16\]](#) [\[19\]](#)
- Explicitly describe:
 - Font category: "bold sans-serif", "vintage script", "graffiti tag", "blackletter".
 - Layout: "centered badge logo", "curved text along top arc", "stacked lines of text", "horizontal banner".
 - Treatment: "solid fill with white outline", "neon glow", "distressed print texture", "3D chrome lettering".
- Add **negative prompts / constraints**: "no extra words, no misspellings, no background text, legible typography".[\[19\]](#) [\[16\]](#)
- For maximum consistency, iterate with image-to-image: generate a layout → re-prompt using that image as reference to refine spelling and spacing (Ideogram, Flux Kontext, Imagen, and

GPT-Image-1 all support image-conditioning flows through their respective APIs).[\[24\]](#) [\[25\]](#) [\[23\]](#)
[\[7\]](#) [\[4\]](#)

Maximum reliable text length (practical):

- Vendor docs avoid hard limits, but hands-on reviews for Ideogram, Midjourney v7, and Imagen 3/4 converge on:
 - **Up to ~4–6 words:** high accuracy across all top models.[\[20\]](#) [\[23\]](#) [\[6\]](#) [\[16\]](#)
 - **Short phrases / 2–3 lines of text (≈20–30 chars each):** good on Ideogram 2/3 and Imagen 4, acceptable on Midjourney v7 and GPT-Image-1, but occasional spacing or minor spelling errors.[\[8\]](#) [\[17\]](#) [\[2\]](#) [\[7\]](#)
 - **Paragraphs:** still unreliable; guidance is to avoid paragraph-length copy in the image and instead compose in vector tools (Figma/Illustrator) after generating the visual frame.[\[26\]](#) [\[23\]](#) [\[6\]](#)

4. Optimal prompt structure (2025 best practice)

Recent 2025 prompting guides across Flux, Ideogram and Google Imagen are fairly aligned on structure.[\[27\]](#) [\[25\]](#) [\[23\]](#) [\[16\]](#)

Recommended structure:

1. **Subject & core task** – “minimalist t-shirt graphic,” “retro 80s poster,” “flat vector logo.”
2. **Exact text in quotes** – “SUNSET SURF CLUB”.
3. **Typography spec** – font category, weight, casing, layout, effects.
4. **Style & medium** – “flat vector illustration,” “screen-print style,” “distressed vintage.”
5. **Color & background** – palette, background color or transparency requirement.
6. **Framing & constraints** – “centered composition,” “no photo background,” “no extra text or watermarks.”

Elements that MUST be included for consistent results:

- A clear **medium** (vector/logo/graphic vs photoreal scene).
- **Aspect ratio / orientation** where supported (e.g. square for chest prints).
- **Foreground/background instructions**, especially for POD (transparent or solid color, no gradients if you want easy printing).[\[25\]](#) [\[3\]](#) [\[27\]](#) [\[6\]](#)

What to avoid (from 2025 guides):

- Over-stuffed prompts with conflicting art directions (e.g. “minimalist flat vector” + “hyper-realistic 3D render”).[\[27\]](#) [\[25\]](#)
- Vague text instructions like “cool logo” without the actual text string.
- Long paragraphs of copy; models interpret text as design elements, not as semantic long-form text.[\[6\]](#) [\[26\]](#)
- Overly specific font names; using categories works better and is more portable.[\[16\]](#) [\[19\]](#)

Style / quality parameters that matter:

- Midjourney: `--stylize` and `--style raw` significantly affect legibility vs artistry; lower `--stylize` improves text clarity. [28] [20] [19]
- Flux / Flux.2: guidance parameters (CFG equivalents) and “Kontext” conditioning provide better prompt adherence and typography control. [29] [24] [1] [27]
- Vertex Imagen 4: guidance scale and safety settings materially change sharpness/text clarity, per Google’s dev guide. [2] [23]

5. T-shirt / merch design specifics

Prompt techniques for POD graphics:

- Start with the medium:
 - “centered t-shirt vector graphic,” “single-color screen-print style,” “front chest print only, no model.” [26] [6]
- Add print constraints:
 - “no photo background,” “no gradients, only flat colors,” “solid outlines suitable for screen printing,” “no tiny details.”
- Define mockup vs artwork:
 - If you only want the **art**, explicitly say “isolated design on plain background, not on a shirt mockup.”

Transparent or solid backgrounds:

- Many web UIs show a checkerboard but the **raw API output is usually opaque**; the common workflow is:
 - Generate on a **flat solid background** (white or #000000), then run automatic background removal (e.g., your own diffusion-masking, `remove.bg` API, or Photoshop/Illustrator scripts). [6] [26]
- In prompts, specify: “plain solid black background,” “no shadows,” “no gradients,” which makes cut-out much easier and more reliable. [25] [27]

Recommended size / aspect ratio for apparel:

- POD platforms typically recommend printing areas around **4500×5400 px @ 300 DPI** or similar; image models rarely natively generate at that resolution, so 2025 practice is:
 - Generate base art at **1024–2048 px square**, then upscale with dedicated upscalers or vectorization (Recraft, Illustrator image trace, etc.). [30] [26] [6]
- For chest-print graphics, **1:1 or 4:5** is usually best to avoid weird cropping when you place onto the print template. [31] [23]

Known print-readiness issues:

- **Fine details, gradients, and noisy textures** often disappear or band when converted to print-ready CMYK; reviewers recommend simplifying colors and shapes in the prompt and

then doing final color-space conversion manually.[\[26\]](#) [\[6\]](#)

- **Text edges** may be slightly fuzzy; vectorizing or re-setting the text in a design tool is still common for genuinely production-ready logos/slogans.[\[7\]](#) [\[6\]](#)

6. Model comparison for text-heavy commercial designs

Leading 2025 models (text-heavy use)

Provider / model	Text accuracy & control (2025)	Notable strengths for merch / text	Caveats
Ideogram 2.0 / 3.0	Among the best; designed for long, stylized typography and "Design" style tuned for text. [7] [5] [8] [16]	Excellent for posters, greeting cards, POD, logos; strong style controls and prompt guide for typography. [7] [8] [16]	Closed source; API access and rate limits vary; still some misspells on very long text.
Midjourney v7	High accuracy for short to medium text; v7 adds dedicated text-rendering capabilities for pro use. [6] [20] [28]	Artistic, high-impact imagery; great for text as part of compositions; Discord and web UI are mature. [6] [26]	API is limited/indirect; fine-grained layout control is trickier than Ideogram.
Google Imagen 4 (Vertex)	Google markets "outstanding text rendering and prompt adherence" and strong enterprise control. [2] [23] [22]	Good for branded creatives, posters, and enterprise workflows via Vertex AI (logging, governance). [2] [23]	GCP/Vertex integration overhead; some lingering issues with complex typography. [22] [31]
GPT-Image-1 (OpenAI)	Good text blocks, especially in diagrams/infographics; slightly weaker than Ideogram/MJ for decorative scripts. [17] [4] [6]	Tight integration with GPT-4o/5.x, assistants, and your existing OpenAI stack; strong safety and watermarking. [4] [11] [14]	No explicit font controls; needs LLM pre-prompting for best layouts.
Flux.2 (Black Forest Labs)	Flux.2 series claims improved typography and prompt understanding vs Flux.1; early tests rate it strong but slightly behind Ideogram for text. [29] [1] [3] [27]	Open-weight options (Klein), good composition, fast inference and strong control via Kontext / depth / canny tools. [29] [24] [1] [27]	Needs more custom workflow (ComfyUI, custom servers); font style is descriptive, not precise.
Adobe Firefly (Creative Cloud)	Very good accuracy for straightforward English phrases; Firefly's killer feature is editable text layers in Photoshop/Illustrator. [6]	Best when you want to tweak text post-generation; integrates perfectly with designer workflows. [6]	Not the most creative or open; licensing and usage constraints to check for merch.
Others (Recraft, Canva, etc.)	Recraft for vector-style art; Canva AI text is modest but compensated by manual text tools. [6] [18]	Great for fast production flows and non-technical teams.	Less control via raw API; may not match your need for full automation.

Best typography control today:

- For **fully automatic, text-in-image layouts**, **Ideogram 2.x/3.0** is currently the most typography-focused and consistent, with Ideogram's own docs emphasizing pro-grade fonts and layout control.[\[8\]](#) [\[7\]](#) [\[16\]](#)
- For **art+text compositions**, **Midjourney v7** and **Imagen 4** are close contenders, especially when you can afford some manual curation.[\[23\]](#) [\[2\]](#) [\[6\]](#) [\[26\]](#)

7. API implementation strategy for your merch generator

Use an LLM to pre-process prompts?

- For a SaaS design generator, 2025 best practice is **absolutely yes**:
 - Parse user intent (theme, niche, mood, text, color constraints, print method) with an LLM.
 - Emit a **structured prompt object**: { text, font_category, layout, style, background, negative_constraints }.
 - Render that into a model-specific natural-language prompt for GPT-Image-1, Ideogram, Midjourney, etc.[\[23\]](#) [\[4\]](#) [\[16\]](#)

Prompt caching / optimization techniques:

- Cache at the **structured representation level**, not just raw strings, so you can:
 - Reuse prompts for A/B model tests (e.g., GPT-Image-1 vs Ideogram).
 - Quickly regenerate variants by only tweaking the typography node or background color node.
- For high-traffic SaaS, store: prompt hash → image IDs + rating/engagement, and train a small re-ranking model or heuristic to bias toward prompts that historically printed well.[\[18\]](#) [\[6\]](#)

Recommended workflow for consistent, high-quality merch designs:

1. User input → LLM normalization

- Use GPT-4o/5.x (or your chosen LLM) to:
 - Validate the slogan (length, banned words).
 - Normalize style tags (e.g., "grunge", "cute kawaii", "retro 80s").

2. Structured design spec

- LLM returns JSON like:
 - text: "NEON TIGER"
 - font_category: "bold condensed sans-serif"
 - layout: "centered badge, text on single line"
 - style: "neon cyberpunk, flat vector screen-print"
 - background: "solid black"
 - negative: "no extra text, no model, no photo"

3. Model-specific prompt renderers

- For GPT-Image-1: single descriptive prompt string; specify size (e.g. square), and mention “t-shirt vector graphic.”^[4]
- For Ideogram: follow its typography prompt guidelines (quotes, specify “Design” style, layout words).^{[7] [16]}
- For Midjourney: add quotes around text, and tune `--stylize / --style raw` for readability.^{[20] [19]}

4. Generation + post-processing

- Generate at 1024–2048 px square.
- Run automatic **background cleanup** and **upscale/vectorization**.
- Optionally re-inject exact text manually or via design templates if absolute spelling precision is required.

5. Human-in-the-loop and telemetry

- Track which model/prompt combos get used or purchased most; feed that data to your LLM prompt-engineer agent to iteratively refine styles and constraints.^{[18] [6]}

This setup gives you: OpenAI for orchestration and generalist image gen, a typography-focused model (Ideogram / MJ / Imagen 4) where text precision matters most, and a clear path to automate yet still keep merch printing constraints under control.

**

1. <https://blogs.nvidia.com/blog/rtx-ai-garage-flux-2-comfyui/>
2. <https://cloud.google.com/blog/products/ai-machine-learning/announcing-veo-3-imagen-4-and-lyria-2-on-vertex-ai>
3. <https://flux-ai.io>
4. <https://platform.openai.com/docs/models/gpt-image-1>
5. <https://www.basedlabs.ai/tools/ideogram-2>
6. <https://ampifire.com/blog/best-ai-image-generators-with-accurate-text-in-2025-reviews-price-free-options/>
7. <https://about.ideogram.ai/2.0>
8. <https://blog.laprompt.com/ai-news/ideogram-3-0-deep-review-comparison-with-2-0-and-2a>
9. <https://platform.openai.com/docs/models>
10. <https://community.openai.com/t/new-gpt-image-model-in-the-api/1239462?page=2>
11. <https://thisisgamethailand.com/technology/openai-image-generator-api/>
12. <https://community.openai.com/t/result-tracking-for-gpt-image-1/1273766>
13. <https://www.cometapi.com/en/gpt-image-1-mini-api/>
14. <https://zapier.com/blog/openai-models/>
15. <https://www.testingcatalog.com/openai-testing-new-image-2-models-on-lm-arena/>
16. <https://docs.ideogram.ai/using-ideogram/prompting-guide/2-prompting-fundamentals/text-and-typography>

17. <https://www.cometapi.com/imagen-3-vs-gpt-image-1-what-is-differences/>
18. <https://ltx.studio/blog/best-ai-image-generator>
19. <https://aichronicler.com/how-to-generate-text-in-midjourney/>
20. <https://www.aimodelsrank.com/reviews/midjourney>
21. <https://fromtexttoimage.com/midjourney-v6-review-2025/>
22. <https://www.imagine.art/blogs/google-imagen-3-features>
23. <https://cloud.google.com/blog/products/ai-machine-learning/a-developers-guide-to-imagen-3-on-vertex-ai>
24. <https://huggingface.co/black-forest-labs/FLUX.1-Kontext-dev>
25. <https://www.promptus.ai/blog/flux-1-kontext-ai-image-editor-complete-guide-2025>
26. <https://www.lovart.ai/blog/ai-illustration-tools-review>
27. <https://docs.comfy.org/tutorials/flux/flux-1-text-to-image>
28. <https://midjourneyv6.org>
29. [https://en.wikipedia.org/wiki/Flux_\(text-to-image_model\)](https://en.wikipedia.org/wiki/Flux_(text-to-image_model))
30. <https://www.cnet.com/tech/services-and-software/best-ai-image-generators/>
31. [https://en.wikipedia.org/wiki/Imagen_\(text-to-image_model\)](https://en.wikipedia.org/wiki/Imagen_(text-to-image_model))
32. <https://bfl.ai>
33. <https://www.siliconflow.com/articles/en/best-open-source-image-generation-models-2025>
34. https://helpx.adobe.com/th_th/firefly/mobile/generate-images-with-text-to-image/generate-images-with-partner-models/flux.html
35. <https://github.com/black-forest-labs/flux>
36. https://www.reddit.com/r/midjourney/comments/18p1jwp/midjourney_v6_can_now_do_text/
37. <https://www.fusionsol.com/blog/gpt-image-1/>
38. <https://www.krea.ai/articles/ai-image-generators-2025>
39. <https://skywork.ai/blog/ideogram-ai-image-generator-vs-linguistic-sign/>
40. <https://openai.com/index/introducing-gpt-5-2/>